

Programming in R and Python

Lecture 10: Linear algebra in Python

Adam Gudyś
Silesian University of Technology
2020

Linear algebra libraries

BLAS (Basic Linear Algebra Subprograms):

- low-level operations (vector-vector, vector-matrix, matrix-matrix),

LAPACK (Linear Algebra Package):

- high-level operations (eigendecomposition; SVD; LU, QR, Cholesky factorizations, etc.).

- Originally implemented in FORTRAN (that time faster than C),
- R – built in the core; Python – NumPy + SciPy packages,
- Possibility to use different implementations:
Intel Math Kernel Library, cuBLAS, OpenBLAS, etc.

Python – installing packages

pip - recommended Python package manager, automatically installed with Python 2 \geq 2.7.9 and Python 3 \geq 3.4.

Where to find it?

- Linux: /usr/bin/ (added to PATH by default),
- Windows: <PythonDir>/scripts/

Pip package manager

Operation	Command
Install latest version	<code>pip install <name></code>
Install specified version	<code>pip install <name>=<version></code>
Install from sources	<code>pip install <src></code>
Upgrade	<code>pip install --upgrade <name></code>
Uninstall	<code>pip uninstall <name></code>

Example:

```
pip install numpy  
pip install scipy
```

Code...

Memory organization

AIDA64 Cache & Memory Benchmark				
	Read	Write	Copy	Latency
Memory	23701 MB/s	24959 MB/s	23728 MB/s	69.2 ns
L1 Cache	762.87 GB/s	390.00 GB/s	744.74 GB/s	1.2 ns
L2 Cache	187.69 GB/s	102.94 GB/s	134.43 GB/s	3.6 ns
L3 Cache	149.70 GB/s	111.54 GB/s	119.71 GB/s	13.2 ns
CPU Type	Mobile QuadCore Intel Core i7-4700MQ (Haswell-MB, rPGA946)			
CPU Stepping	C0			
CPU Clock	3192.6 MHz (original: 2400 MHz, overclock: 33%)			
CPU FSB	99.8 MHz (original: 100 MHz)			
CPU Multiplier	32x	North Bridge Clock		3192.6 MHz

32GB

4x32KB

4x256KB

6MB

Experiment

Matrix multiplication: $C_{M \times N} = A_{M \times K} * B_{K \times N}$

- code written in C, matrices stored as 1D arrays of double.

```
for (int i = 0; i < M; ++i)
    for (int j = 0; j < N; ++j)
        for (int k = 0; k < K; ++k)
            C[i*N+j] += A[i*K+k] + B[k*N+j];
```

Experimental setting:

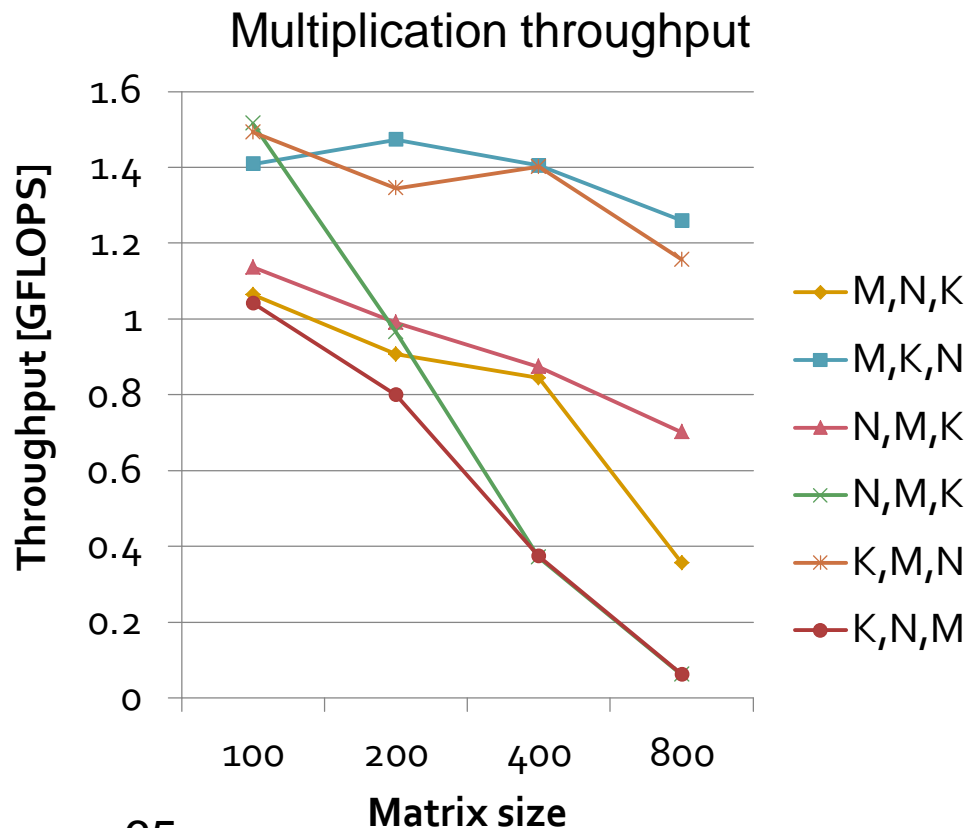
- Matrices were assumed to be square: $M = N = K$.
- Different orderings of M, N, K loops were analyzed:
 - each time, same number of operations performed!

Results

Execution times for different loop orderings (milliseconds).

size	100	200	400	800
M, N, K	0.9	8.8	75.8	1437.7
M, K, N	0.7	5.4	45.6	406.5
N, M, K	0.9	8.1	73.3	730.5
N, M, K	0.7	8.3	172.7	8366.2
K, M, N	0.7	6.0	45.7	442.8
K, N, M	1.0	10.0	171.2	8203.0

i7-4700MQ double precision performance: ~65 GFLOPS



Thank you for your attention!